

```
In [1]: import tensorflow as tf  
import json  
import numpy as np  
from matplotlib import pyplot as plt
```

```
In [2]: images = tf.data.Dataset.list_files('D:\\\\images_2\\\\images\\\\*.jpg')
```

```
In [3]: images.as_numpy_iterator().next()
```

```
Out[3]: b'D:\\\\images_2\\\\images\\\\006.jpg'
```

```
In [4]: def load_image(x):  
    byte_img = tf.io.read_file(x)  
    img = tf.io.decode_jpeg(byte_img)  
    return img
```

```
In [5]: images = images.map(load_image)
```

```
In [6]: images.as_numpy_iterator().next()
```

```
Out[6]: array([[[180, 180, 190],
   [179, 182, 191],
   [179, 183, 194],
   ...,
   [159, 173, 173],
   [163, 177, 178],
   [166, 180, 181]],

   [[180, 180, 188],
   [177, 180, 187],
   [174, 181, 189],
   ...,
   [166, 180, 180],
   [167, 181, 182],
   [167, 181, 182]],

   [[179, 182, 187],
   [176, 181, 185],
   [173, 180, 186],
   ...,
   [170, 184, 184],
   [169, 183, 183],
   [164, 178, 178]],

   ...,

   [[105, 90, 67],
   [105, 87, 63],
   [107, 85, 61],
   ...,
   [135, 99, 67],
   [127, 91, 57],
   [138, 104, 69]],

   [[118, 98, 71],
   [113, 92, 65],
   [113, 87, 62],
   ...,
   [137, 101, 69],
   [125, 91, 56],
   [131, 99, 61]],

   [[121, 100, 71],
   [113, 89, 61],
   [113, 85, 61],
   ...,
   [138, 104, 69],
   [128, 94, 59],
   [128, 96, 58]]], dtype=uint8)
```

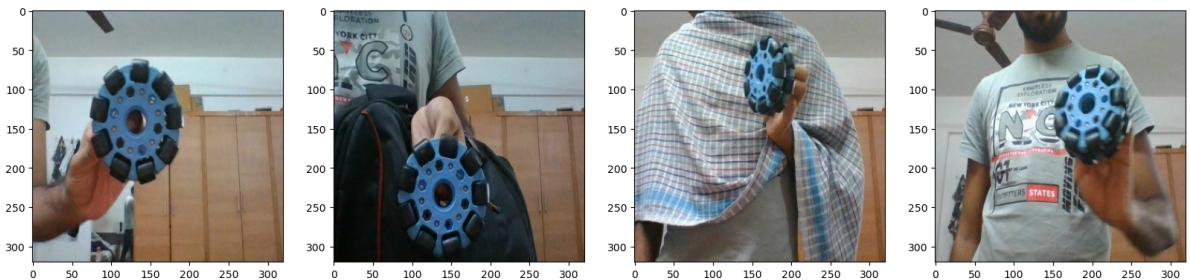
```
In [7]: type(images)
```

```
Out[7]: tensorflow.python.data.ops.map_op._MapDataset
```

```
In [8]: image_generator = images.batch(4).as_numpy_iterator()
```

```
In [9]: plot_images = image_generator.next()
```

```
In [10]: fig, ax = plt.subplots(ncols = 4, figsize = (20,20))
for idx, image in enumerate(plot_images):
    ax[idx].imshow(image)
plt.show()
```



In [11]: `50*.7`

Out[11]: `35.0`

In [12]: `50*.15`

Out[12]: `7.5`

```
In [13]: import os
for folder in ['train', 'test', 'val']:
    for file in os.listdir(os.path.join('D:\\images_2', folder, 'images')):
        filename = file.split('.')[0] + '.json'
        existing_filepath = os.path.join('D:\\images_2', 'labels', filename)
        if os.path.exists(existing_filepath):
            new_filepath = os.path.join('D:\\images_2', folder, 'labels', filename)
            os.replace(existing_filepath, new_filepath)
```

In [14]: `import albumentations as alb`

```
In [15]: augmentor = alb.Compose([alb.RandomCrop(width=320, height=320),
                             alb.HorizontalFlip(p=0.5),
                             alb.RandomBrightnessContrast(p=0.2),
                             alb.RandomGamma(p=0.2),
                             alb.RGBShift(p=0.2),
                             alb.VerticalFlip(p=0.5)],
                             bbox_params=alb.BboxParams(format='albumentations',
                             label_fields=['class_labels']))
```

```
In [16]: import cv2
img = cv2.imread(os.path.join('D:\\images_2', 'train', 'images', '002.jpg'))
```

In [17]: `img`

```
Out[17]: array([[[202, 189, 187],
   [204, 190, 184],
   [209, 193, 181],
   ...,
   [194, 179, 176],
   [197, 176, 178],
   [197, 176, 178]],

   [[205, 193, 189],
   [209, 196, 188],
   [213, 197, 184],
   ...,
   [195, 180, 177],
   [196, 178, 177],
   [198, 177, 179]],

   [[205, 196, 187],
   [208, 198, 188],
   [218, 205, 189],
   ...,
   [196, 181, 178],
   [197, 179, 178],
   [197, 179, 178]],

   ...,

   [[ 13,  11,  10],
   [ 13,  14,  12],
   [ 14,  16,  17],
   ...,
   [ 69,  93, 139],
   [ 61,  85, 131],
   [ 77, 102, 146]],

   [[ 17,  13,  12],
   [ 17,  15,  14],
   [ 17,  16,  18],
   ...,
   [ 71,  97, 143],
   [ 65,  91, 137],
   [ 72,  99, 143]],

   [[ 17,  13,  12],
   [ 16,  14,  13],
   [ 19,  16,  18],
   ...,
   [ 74, 100, 146],
   [ 67,  93, 139],
   [ 65,  92, 136]]], dtype=uint8)
```

```
In [18]: with open(os.path.join('D:\\images_2', 'train', 'labels', '002.json'), 'r') as f:
    label = json.load(f)
```

```
In [19]: label['shapes'][0]['points']
```

```
Out[19]: [[29.965277777777757, 130.48611111111111], [236.15234375, 319.0]]
```

```
In [20]: coords = [0,0,0,0]
coords[0] = label['shapes'][0]['points'][0][0]
coords[1] = label['shapes'][0]['points'][0][1]
coords[2] = label['shapes'][0]['points'][1][0]
coords[3] = label['shapes'][0]['points'][1][1]
```

```
In [21]: coords
```

```
Out[21]: [29.965277777777757, 130.48611111111111, 236.15234375, 319.0]
```

```
In [22]: coords = list(np.divide(coords,[320,320,320,320]))
```

```
In [23]: coords
```

```
Out[23]: [0.0936414930555555, 0.4077690972222222, 0.73797607421875, 0.996875]
```

```
In [24]: augmented = augmentor(image=img, bboxes=[coords], class_labels=['face'])
```

```
In [25]: augmented.keys()
```

```
Out[25]: dict_keys(['image', 'bboxes', 'class_labels'])
```

```
In [26]: augmented['bboxes'][0][2:]
```

```
Out[26]: (0.9063585069444445, 0.5922309027777778)
```

```
In [27]: augmented['bboxes']
```

```
Out[27]: [(0.26202392578125,  
          0.0031250000000000444,  
          0.9063585069444445,  
          0.5922309027777778)]
```

```
In [28]: augmented['image']
```

```
Out[28]: array([[[ 65,  92, 136],
   [ 67,  93, 139],
   [ 74, 100, 146],
   ...,
   [ 19,  16,  18],
   [ 16,  14,  13],
   [ 17,  13,  12]],

   [[ 72,  99, 143],
   [ 65,  91, 137],
   [ 71,  97, 143],
   ...,
   [ 17,  16,  18],
   [ 17,  15,  14],
   [ 17,  13,  12]],

   [[ 77, 102, 146],
   [ 61,  85, 131],
   [ 69,  93, 139],
   ...,
   [ 14,  16,  17],
   [ 13,  14,  12],
   [ 13,  11,  10]],

   ...,

   [[197, 179, 178],
   [197, 179, 178],
   [196, 181, 178],
   ...,
   [218, 205, 189],
   [208, 198, 188],
   [205, 196, 187]],

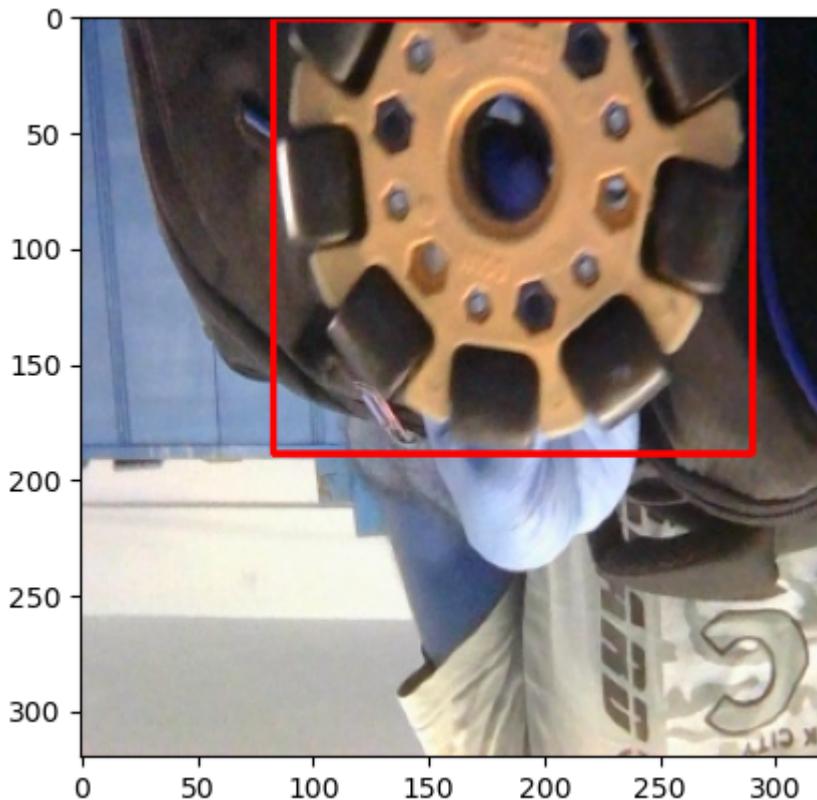
   [[198, 177, 179],
   [196, 178, 177],
   [195, 180, 177],
   ...,
   [213, 197, 184],
   [209, 196, 188],
   [205, 193, 189]],

   [[197, 176, 178],
   [197, 176, 178],
   [194, 179, 176],
   ...,
   [209, 193, 181],
   [204, 190, 184],
   [202, 189, 187]]], dtype=uint8)
```

```
In [29]: cv2.rectangle(augmented['image'],
                     tuple(np.multiply(augmented['bboxes'][0][:2], [320,320]).astype(int)),
                     tuple(np.multiply(augmented['bboxes'][0][2:], [320,320]).astype(int)),
                     (255,0,0), 2)

plt.imshow(augmented['image'])
```

```
Out[29]: <matplotlib.image.AxesImage at 0x17356a8b370>
```



```
In [30]: for partition in ['train', 'test', 'val']:
    for image in os.listdir(os.path.join('D:\\images_2', partition, 'images')):
        img = cv2.imread(os.path.join('D:\\images_2', partition, 'images', image))

        coords = [0,0,0.00001,0.00001]
        label_path = os.path.join('D:\\images_2', partition, 'labels', f'{image}.json')
        if os.path.exists(label_path):
            with open(label_path, 'r') as f:
                label = json.load(f)

                coords[0] = label['shapes'][0]['points'][0][0]
                coords[1] = label['shapes'][0]['points'][0][1]
                coords[2] = label['shapes'][0]['points'][1][0]
                coords[3] = label['shapes'][0]['points'][1][1]
                coords = list(np.divide(coords, [320,320,320,320]))

        try:
            for x in range(34):
                augmented = augmentor(image=img, bboxes=[coords], class_labels=['fan'])
                cv2.imwrite(os.path.join('D:\\aug_data', partition, 'images', f'{image}_{x}.jpg'))

                annotation = {}
                annotation['image'] = image

                if os.path.exists(label_path):
                    if len(augmented['bboxes']) == 0:
                        annotation['bbox'] = [0,0,0,0]
                        annotation['class'] = 0
                    else:
                        annotation['bbox'] = augmented['bboxes'][0]
                        annotation['class'] = 1
                else:
                    annotation['bbox'] = [0,0,0,0]
                    annotation['class'] = 0

                with open(os.path.join('D:\\aug_data', partition, 'labels', f'{image}_{x}.json'), 'w') as f:
                    json.dump(annotation, f)
```

```
        json.dump(annotation, f)
```

```
    except Exception as e:
        print(e)
```

```
In [31]: train_images = tf.data.Dataset.list_files('D:\\\\aug_data\\\\train\\\\images\\\\*.jpg', shuffle=True)
train_images = train_images.map(load_image)
train_images = train_images.map(lambda x: tf.image.resize(x, (120,120)))
train_images = train_images.map(lambda x: x/255)
```

```
In [32]: test_images = tf.data.Dataset.list_files('D:\\\\aug_data\\\\test\\\\images\\\\*.jpg', shuffle=True)
test_images = test_images.map(load_image)
test_images = test_images.map(lambda x: tf.image.resize(x, (120,120)))
test_images = test_images.map(lambda x: x/255)
```

```
In [33]: val_images = tf.data.Dataset.list_files('D:\\\\aug_data\\\\val\\\\images\\\\*.jpg', shuffle=True)
val_images = val_images.map(load_image)
val_images = val_images.map(lambda x: tf.image.resize(x, (120,120)))
val_images = val_images.map(lambda x: x/255)
```

```
In [34]: def load_labels(label_path):
    with open(label_path.numpy(), 'r', encoding = "utf-8") as f:
        label = json.load(f)

    return [label['class']], label['bbox']
```

```
In [35]: train_labels = tf.data.Dataset.list_files('D:\\\\aug_data\\\\train\\\\labels\\\\*.json', shuffle=True)
train_labels = train_labels.map(lambda x: tf.py_function(load_labels, [x], [tf.uint8]))
```

```
In [36]: test_labels = tf.data.Dataset.list_files('D:\\\\aug_data\\\\test\\\\labels\\\\*.json', shuffle=True)
test_labels = test_labels.map(lambda x: tf.py_function(load_labels, [x], [tf.uint8]))
```

```
In [37]: val_labels = tf.data.Dataset.list_files('D:\\\\aug_data\\\\val\\\\labels\\\\*.json', shuffle=True)
val_labels = val_labels.map(lambda x: tf.py_function(load_labels, [x], [tf.uint8]))
```

```
In [38]: train_labels.as_numpy_iterator().next()
```

```
Out[38]: (array([1], dtype=uint8),
          array([0.262    , 0.003124, 0.9062   , 0.5923   ], dtype=float16))
```

```
In [39]: len(train_images), len(train_labels), len(test_images), len(test_labels), len(val_images)
```

```
Out[39]: (1190, 1190, 420, 420, 272, 272)
```

```
In [40]: train = tf.data.Dataset.zip((train_images, train_labels))
train = train.shuffle(3000)
train = train.batch(4)
train = train.prefetch(4)
```

```
In [41]: test = tf.data.Dataset.zip((test_images, test_labels))
test = test.shuffle(1300)
test = test.batch(4)
test = test.prefetch(4)
```

```
In [42]: val = tf.data.Dataset.zip((val_images, val_labels))
val = val.shuffle(1000)
val = val.batch(4)
val = val.prefetch(4)
```

```
In [43]: train.as_numpy_iterator().next()[1]
```

```
Out[43]: (array([[1],
       [1],
       [1],
       [1]], dtype=uint8),
 array([[0.3594 , 0.3003 , 0.814  , 0.8896 ],
       [0.01009, 0.00846, 0.495  , 0.6543 ],
       [0.4702 , 0.3513 , 0.6787 , 0.676  ],
       [0.3345 , 0.4675 , 0.5625 , 0.767  ]], dtype=float16))
```

```
In [44]: data_samples = train.as_numpy_iterator()
```

```
In [45]: res = data_samples.next()
```

```
In [46]: fig, ax = plt.subplots(ncols=4, figsize=(20,20))
for idx in range(4):
    sample_image = res[0][idx]
    sample_coords = res[1][1][idx]

    cv2.rectangle(sample_image,
                  tuple(np.multiply(sample_coords[:2], [120,120]).astype(int)),
                  tuple(np.multiply(sample_coords[2:], [120,120]).astype(int)),
                  (255,0,0), 2)

    ax[idx].imshow(sample_image)
```

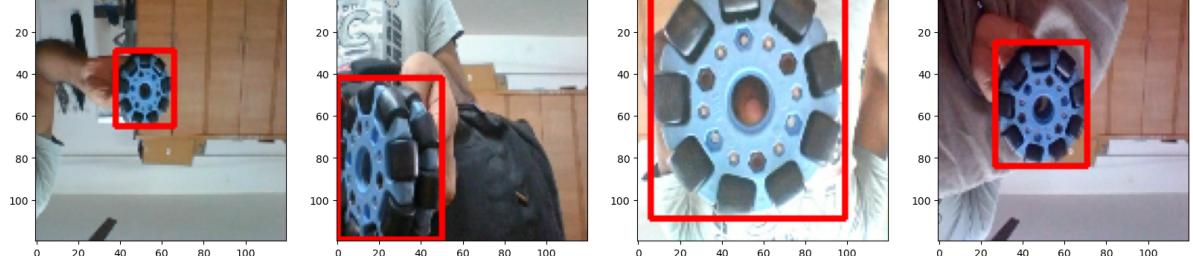
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).



```
In [47]: from tensorflow.keras.models import Model
from tensorflow.keras.layers import Input, Conv2D, Dense, GlobalMaxPooling2D
from tensorflow.keras.applications import VGG16
```

```
In [48]: vgg = VGG16(include_top=False)
```

```
In [49]: vgg.summary()
```

Model: "vgg16"

Layer (type)	Output Shape	Param #
<hr/>		
input_1 (InputLayer)	[ (None, None, None, 3) ]	0
block1_conv1 (Conv2D)	(None, None, None, 64)	1792
block1_conv2 (Conv2D)	(None, None, None, 64)	36928
block1_pool (MaxPooling2D)	(None, None, None, 64)	0
block2_conv1 (Conv2D)	(None, None, None, 128)	73856
block2_conv2 (Conv2D)	(None, None, None, 128)	147584
block2_pool (MaxPooling2D)	(None, None, None, 128)	0
block3_conv1 (Conv2D)	(None, None, None, 256)	295168
block3_conv2 (Conv2D)	(None, None, None, 256)	590080
block3_conv3 (Conv2D)	(None, None, None, 256)	590080
block3_pool (MaxPooling2D)	(None, None, None, 256)	0
block4_conv1 (Conv2D)	(None, None, None, 512)	1180160
block4_conv2 (Conv2D)	(None, None, None, 512)	2359808
block4_conv3 (Conv2D)	(None, None, None, 512)	2359808
block4_pool (MaxPooling2D)	(None, None, None, 512)	0
block5_conv1 (Conv2D)	(None, None, None, 512)	2359808
block5_conv2 (Conv2D)	(None, None, None, 512)	2359808
block5_conv3 (Conv2D)	(None, None, None, 512)	2359808
block5_pool (MaxPooling2D)	(None, None, None, 512)	0
<hr/>		
Total params:	14,714,688	
Trainable params:	14,714,688	
Non-trainable params:	0	

In [50]:

```
def build_model():
    input_layer = Input(shape=(120,120,3))

    vgg = VGG16(include_top=False)(input_layer)

    # Classification Model
    f1 = GlobalMaxPooling2D()(vgg)
    class1 = Dense(2048, activation='relu')(f1)
    class2 = Dense(1, activation='sigmoid')(class1)

    # Bounding box model
    f2 = GlobalMaxPooling2D()(vgg)
    regress1 = Dense(2048, activation='relu')(f2)
    regress2 = Dense(4, activation='sigmoid')(regress1)
```

```
wheeltracker = Model(inputs=input_layer, outputs=[class2, regress2])
return wheeltracker
```

In [51]: `wheeltracker = build_model()`

In [52]: `wheeltracker.summary()`

Model: "model"

Layer (type)	Output Shape	Param #	Connected to
<hr/>			
<hr/>			
input_2 (InputLayer)	[(None, 120, 120, 3 0 )]		[]
vgg16 (Functional)	(None, None, None, 512)	14714688	['input_2[0][0]']
global_max_pooling2d (GlobalMa xPooling2D)	(None, 512)	0	['vgg16[0][0]']
global_max_pooling2d_1 (Global MaxPooling2D)	(None, 512)	0	['vgg16[0][0]']
dense (Dense) ing2d[0][0]']	(None, 2048)	1050624	['global_max_pool
dense_2 (Dense) ing2d_1[0][0]']	(None, 2048)	1050624	['global_max_pool
dense_1 (Dense)	(None, 1)	2049	['dense[0][0]']
dense_3 (Dense)	(None, 4)	8196	['dense_2[0][0]']
<hr/>			
<hr/>			
Total params: 16,826,181			
Trainable params: 16,826,181			
Non-trainable params: 0			

In [53]: `X,y = train.as_numpy_iterator().next()`

In [54]: `X.shape`

Out[54]: `(4, 120, 120, 3)`

In [55]: `classes, coords = wheeltracker.predict(X)`

1/1 [=====] - 1s 808ms/step

In [56]: `classes,coords`

Out[56]: `(array([[0.63973665],
 [0.6732118 ],
 [0.5954127 ],
 [0.49985802]], dtype=float32),
 array([[0.3698088 , 0.2248449 , 0.62256306, 0.51001376],
 [0.38658082, 0.2305657 , 0.55798966, 0.42848614],
 [0.3334335 , 0.20294228, 0.5998287 , 0.55069065],
 [0.36772338, 0.27767774, 0.69055504, 0.5154221 ]], dtype=float32))`

```
In [57]: batches_per_epoch = len(train)
lr_decay = (1./0.75 -1)/batches_per_epoch
```

```
In [58]: import tensorflow as tf

initial_learning_rate = 0.0001
lr_decay = 1e-6 # Example decay value

# Define the Learning rate schedule with decay
lr_schedule = tf.keras.optimizers.schedules.ExponentialDecay(
    initial_learning_rate,
    decay_steps=10000, # Adjust decay steps as needed
    decay_rate=0.96,
    staircase=True)

# Create the Adam optimizer using the Learning rate schedule
opt = tf.keras.optimizers.Adam(learning_rate=lr_schedule)
```

```
In [59]: def localization_loss(y_true, yhat):
    delta_coord = tf.reduce_sum(tf.square(y_true[:, :2] - yhat[:, :2]))

    h_true = y_true[:, 3] - y_true[:, 1]
    w_true = y_true[:, 2] - y_true[:, 0]

    h_pred = yhat[:, 3] - yhat[:, 1]
    w_pred = yhat[:, 2] - yhat[:, 0]

    delta_size = tf.reduce_sum(tf.square(w_true - w_pred) + tf.square(h_true-h_pred))

    return delta_coord + delta_size
```

```
In [60]: classloss = tf.keras.losses.BinaryCrossentropy()
regressloss = localization_loss
```

```
In [61]: localization_loss(y[1], coords)
```

```
Out[61]: <tf.Tensor: shape=(), dtype=float32, numpy=0.9014548>
```

```
In [62]: classloss(y[0], classes)
```

```
Out[62]: <tf.Tensor: shape=(), dtype=float32, numpy=0.5135812>
```

```
In [63]: regressloss(y[1], coords)
```

```
Out[63]: <tf.Tensor: shape=(), dtype=float32, numpy=0.9014548>
```

```
In [64]: class WheelTracker(Model):
    def __init__(self, eyetracker, **kwargs):
        super().__init__(**kwargs)
        self.model = eyetracker

    def compile(self, opt, classloss, localizationloss, **kwargs):
        super().compile(**kwargs)
        self.classloss = classloss
        self.localizationloss = localizationloss
        self.opt = opt

    def train_step(self, batch, **kwargs):
        X, y = batch
```

```
with tf.GradientTape() as tape:
    classes, coords = self.model(X, training=True)

    batch_classloss = self.closs(y[0], classes)
    batch_localizationloss = self.lloss(tf.cast(y[1], tf.float32), coords)

    total_loss = batch_localizationloss+0.5*batch_classloss

    grad = tape.gradient(total_loss, self.model.trainable_variables)

opt.apply_gradients(zip(grad, self.model.trainable_variables))

return {"total_loss":total_loss, "class_loss":batch_classloss, "regress_loss":batch_localizationloss}

def test_step(self, batch, **kwargs):
    X, y = batch

    classes, coords = self.model(X, training=False)

    batch_classloss = self.closs(y[0], classes)
    batch_localizationloss = self.lloss(tf.cast(y[1], tf.float32), coords)
    total_loss = batch_localizationloss+0.5*batch_classloss

    return {"total_loss":total_loss, "class_loss":batch_classloss, "regress_loss":batch_localizationloss}

def call(self, X, **kwargs):
    return self.model(X, **kwargs)
```

In [65]: model = WheelTracker(wheeltracker)

In [66]: model.compile(opt, classloss, regressloss)

In [67]: logdir='logs'

In [68]: tensorboard\_callback = tf.keras.callbacks.TensorBoard(log\_dir=logdir)

In [69]: hist = model.fit(train, epochs=10, validation\_data=val, callbacks=[tensorboard\_callback])

```
Epoch 1/10
298/298 [=====] - 272s 887ms/step - total_loss: 0.0905 - class_loss: 0.0047 - regress_loss: 0.0882 - val_total_loss: 0.0854 - val_class_loss: 1.8344e-05 - val_regress_loss: 0.0854
Epoch 2/10
298/298 [=====] - 225s 753ms/step - total_loss: 0.0125 - class_loss: 2.0134e-05 - regress_loss: 0.0125 - val_total_loss: 0.0476 - val_class_loss: 4.0829e-06 - val_regress_loss: 0.0476
Epoch 3/10
298/298 [=====] - 222s 744ms/step - total_loss: 0.0133 - class_loss: 6.0874e-06 - regress_loss: 0.0133 - val_total_loss: 0.0090 - val_class_loss: 1.2368e-06 - val_regress_loss: 0.0090
Epoch 4/10
298/298 [=====] - 219s 732ms/step - total_loss: 0.0062 - class_loss: 2.8463e-06 - regress_loss: 0.0062 - val_total_loss: 0.0105 - val_class_loss: 1.0133e-06 - val_regress_loss: 0.0105
Epoch 5/10
298/298 [=====] - 218s 729ms/step - total_loss: 0.0071 - class_loss: 1.9241e-06 - regress_loss: 0.0071 - val_total_loss: 0.0552 - val_class_loss: 1.8477e-06 - val_regress_loss: 0.0552
Epoch 6/10
298/298 [=====] - 219s 733ms/step - total_loss: 0.0055 - class_loss: 1.3490e-06 - regress_loss: 0.0055 - val_total_loss: 0.0264 - val_class_loss: 1.2517e-06 - val_regress_loss: 0.0264
Epoch 7/10
298/298 [=====] - 221s 739ms/step - total_loss: 0.0065 - class_loss: 8.9168e-07 - regress_loss: 0.0065 - val_total_loss: 0.0123 - val_class_loss: 0.0000e+00 - val_regress_loss: 0.0123
Epoch 8/10
298/298 [=====] - 220s 736ms/step - total_loss: 0.0046 - class_loss: 6.4718e-07 - regress_loss: 0.0046 - val_total_loss: 0.0027 - val_class_loss: 5.9605e-08 - val_regress_loss: 0.0027
Epoch 9/10
298/298 [=====] - 221s 738ms/step - total_loss: 0.0073 - class_loss: 7.0056e-07 - regress_loss: 0.0073 - val_total_loss: 0.0027 - val_class_loss: 5.3644e-07 - val_regress_loss: 0.0027
Epoch 10/10
298/298 [=====] - 222s 743ms/step - total_loss: 0.0068 - class_loss: 7.2388e-07 - regress_loss: 0.0068 - val_total_loss: 0.0073 - val_class_loss: 5.3644e-07 - val_regress_loss: 0.0073
```

In [70]: hist.history

```
Out[70]: {'total_loss': [0.0067139542661607265,
 0.0010398124577477574,
 0.0059084463864564896,
 0.00043239188380539417,
 0.0009732298785820603,
 0.0014478960074484348,
 0.0013017215533182025,
 0.002146644052118063,
 0.001702834153547883,
 0.0014607426710426807],
 'class_loss': [0.00010026506788562983,
 1.2964095731149428e-05,
 6.854538696643431e-07,
 8.046631023717055e-07,
 1.3113038903611596e-06,
 2.8312267659202917e-06,
 0.0,
 2.9802322387695312e-08,
 0.0,
 8.344657089764951e-07],
 'regress_loss': [0.006663821637630463,
 0.0010333304526284337,
 0.005908103659749031,
 0.0004319895524531603,
 0.000972574227489531,
 0.0014464803971350193,
 0.0013017215533182025,
 0.002146629150956869,
 0.001702834153547883,
 0.001460325438529253],
 'val_total_loss': [0.08537231385707855,
 0.04760414734482765,
 0.00899313110858202,
 0.010493303649127483,
 0.055169083178043365,
 0.02640898898243904,
 0.01231568306684494,
 0.0026930654421448708,
 0.0027068983763456345,
 0.007261715363711119],
 'val_class_loss': [1.8343856936553493e-05,
 4.082930445292732e-06,
 1.2367975159577327e-06,
 1.0132800980500178e-06,
 1.8477462617738638e-06,
 1.251698790838418e-06,
 0.0,
 5.960465188081798e-08,
 5.364421440390288e-07,
 5.364420871956099e-07],
 'val_regress_loss': [0.08536314219236374,
 0.047602105885744095,
 0.008992512710392475,
 0.010492797009646893,
 0.05516815930604935,
 0.0264083631336689,
 0.01231568306684494,
 0.002693035639822483,
 0.002706630155444145,
 0.0072614471428096294]}
```

```
In [71]: fig, ax = plt.subplots(ncols=3, figsize=(20,5))
ax[0].plot(hist.history['total_loss'], color='teal', label='loss')
```

```

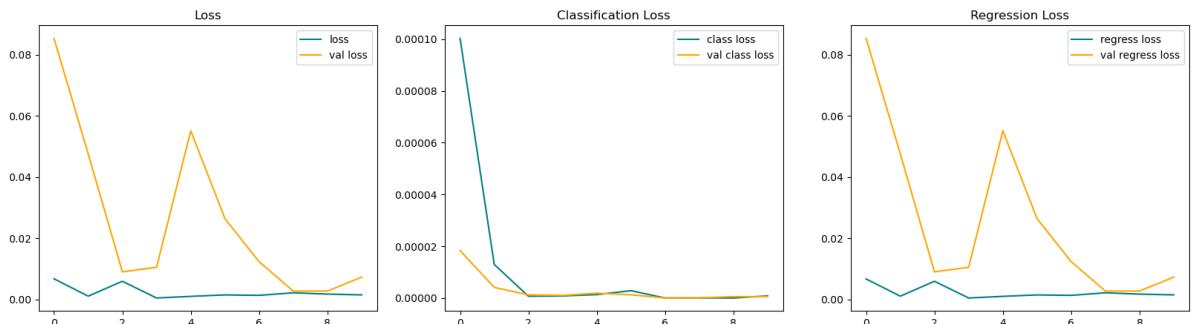
ax[0].plot(hist.history['val_total_loss'], color='orange', label='val loss')
ax[0].title.set_text('Loss')
ax[0].legend()

ax[1].plot(hist.history['class_loss'], color='teal', label='class loss')
ax[1].plot(hist.history['val_class_loss'], color='orange', label='val class loss')
ax[1].title.set_text('Classification Loss')
ax[1].legend()

ax[2].plot(hist.history['regress_loss'], color='teal', label='regress loss')
ax[2].plot(hist.history['val_regress_loss'], color='orange', label='val regress loss')
ax[2].title.set_text('Regression Loss')
ax[2].legend()

plt.show()

```



In [72]: `test_data = test.as_numpy_iterator()`

In [98]: `test_sample = test_data.next()`

In [99]: `yhat = wheeltracker.predict(test_sample[0])`

1/1 [=====] - 0s 100ms/step

```

In [100...]:
fig, ax = plt.subplots(ncols=4, figsize=(20,20))
for idx in range(4):
    sample_image = test_sample[0][idx]
    sample_coords = yhat[1][idx]

    if yhat[0][idx] > 0.9:
        cv2.rectangle(sample_image,
                      tuple(np.multiply(sample_coords[:2], [120,120]).astype(int)),
                      tuple(np.multiply(sample_coords[2:], [120,120]).astype(int)),
                      (255,0,0), 1)

    ax[idx].imshow(sample_image)

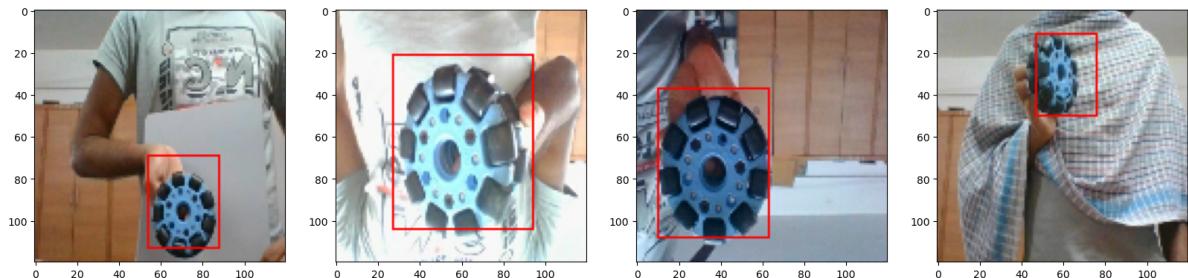
```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).



```
In [92]: from tensorflow.keras.models import load_model
```

```
In [87]: wheeltracker.save('wheeltracker.h5')
```

WARNING:tensorflow:Compiled the loaded model, but the compiled metrics have yet to be built. `model.compile\_metrics` will be empty until you train or evaluate the model.

```
In [88]: wheeltracker = load_model('wheeltracker.h5')
```

WARNING:tensorflow:No training configuration found in the save file, so the model was \*not\* compiled. Compile it manually.

```
In [80]: cap = cv2.VideoCapture(1)
while cap.isOpened():
    _, frame = cap.read()
    frame = frame[50:500, 50:500,:]

    rgb = cv2.cvtColor(frame, cv2.COLOR_BGR2RGB)
    resized = tf.image.resize(rgb, (120,120))

    yhat = wheeltracker.predict(np.expand_dims(resized/255,0))
    sample_coords = yhat[1][0]

    if yhat[0] > 0.5:
        # Controls the main rectangle
        cv2.rectangle(frame,
                      tuple(np.multiply(sample_coords[:2], [450,450]).astype(int)),
                      tuple(np.multiply(sample_coords[2:], [450,450]).astype(int)),
                      (255,0,0), 2)
        # Controls the label rectangle
        cv2.rectangle(frame,
                      tuple(np.add(np.multiply(sample_coords[:2], [450,450]).astype(int),
                                  [0,-30])),
                      tuple(np.add(np.multiply(sample_coords[:2], [450,450]).astype(int),
                                  [80,0])),
                      (255,0,0), -1)

        # Controls the text rendered
        cv2.putText(frame, 'face', tuple(np.add(np.multiply(sample_coords[:2], [450,450]).astype(int),
                                              [0,-5])),
                   cv2.FONT_HERSHEY_SIMPLEX, 1, (255,255,255), 2, cv2.LINE_AA)

    cv2.imshow('EyeTrack', frame)

    if cv2.waitKey(1) & 0xFF == ord('q'):
        break
cap.release()
cv2.destroyAllWindows()
```

```
In [ ]:
```